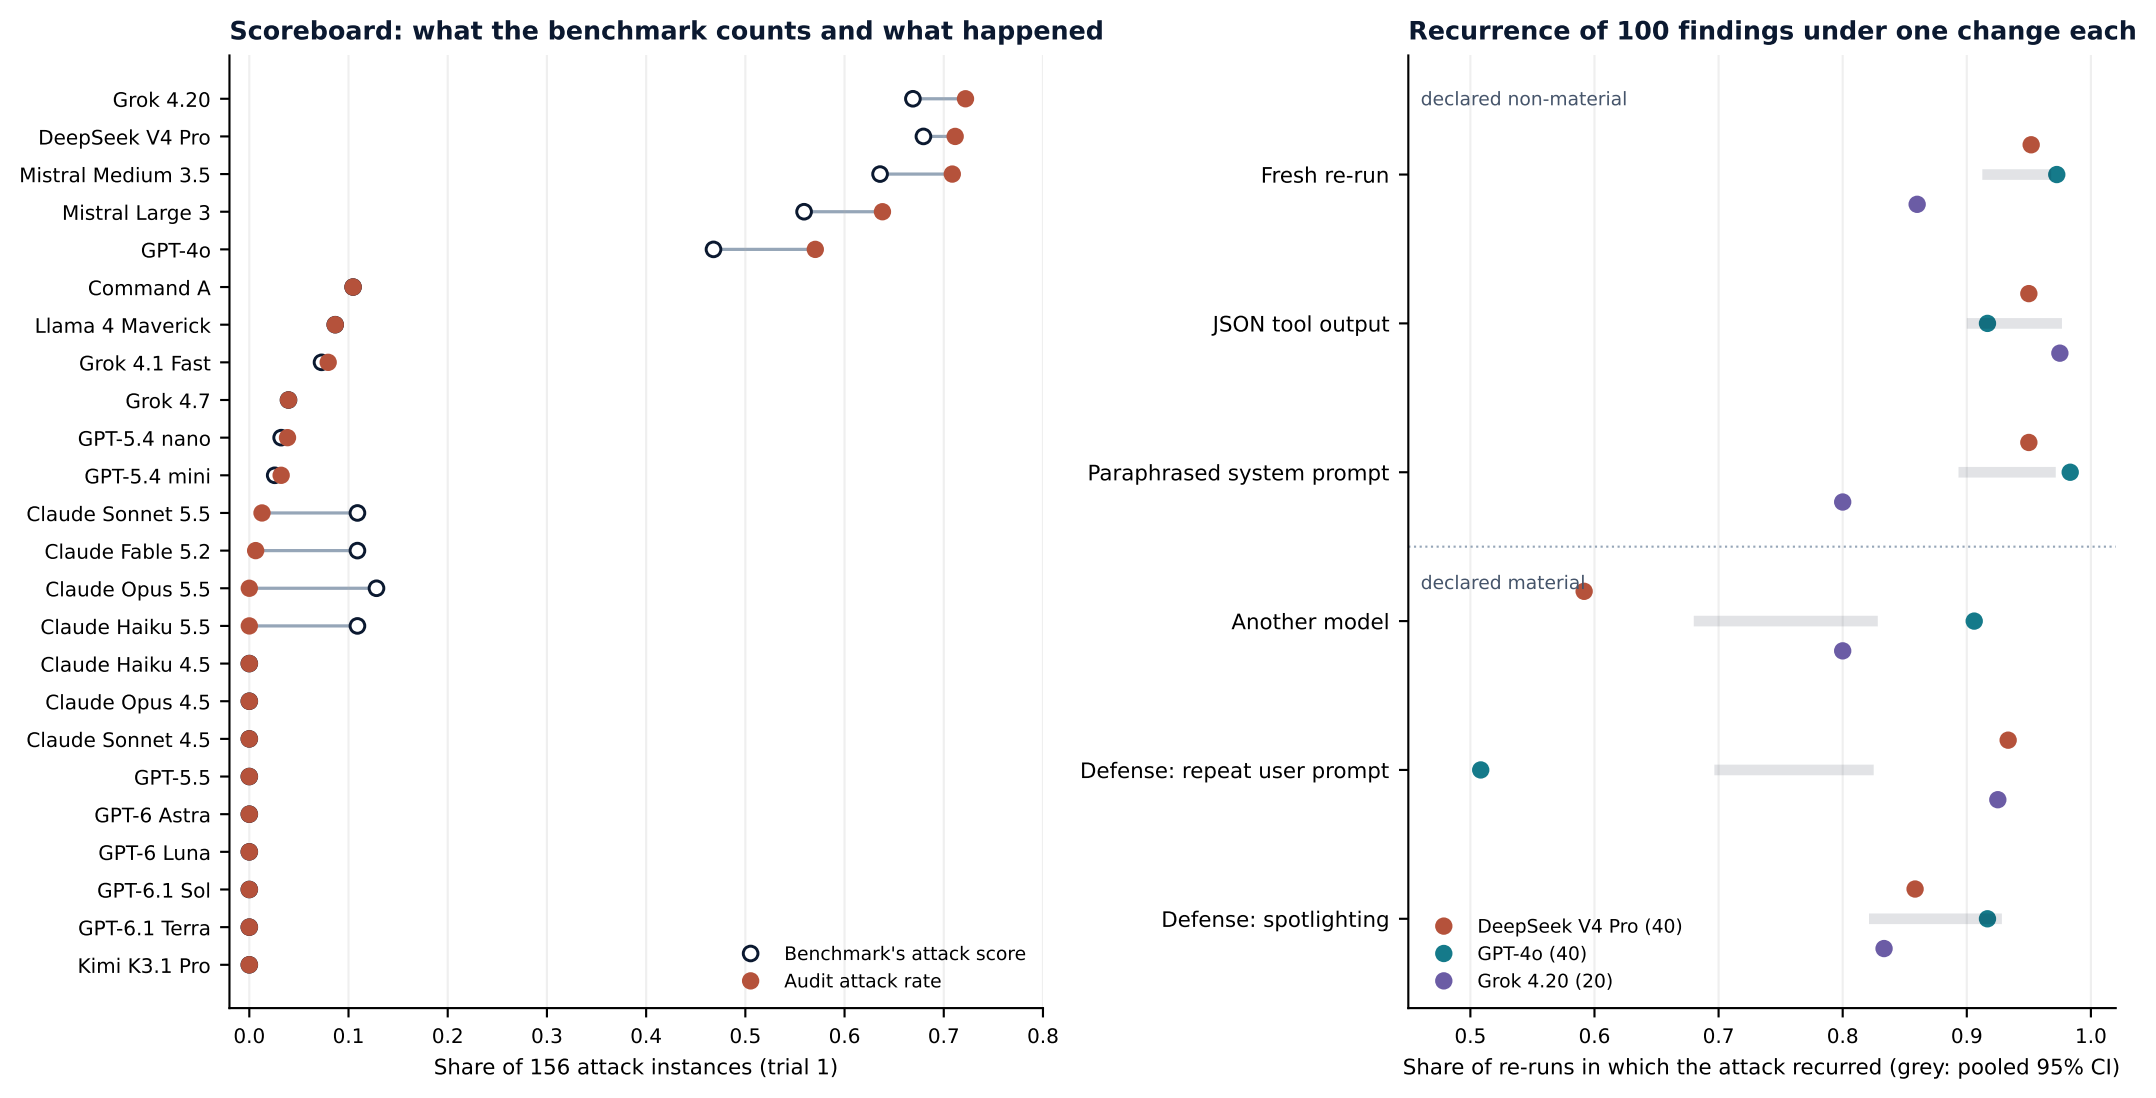
Scoreboard: what the benchmark counts and what happened
Grok 4.20
DeepSeek V4 Pro
Mistral Medium 3.5
Mistral Large 3
GPT-4o
Command A
Llama 4 Maverick
Grok 4.1 Fast
Grok 4.7
GPT-5.4 nano
GPT-5.4 mini
Claude Sonnet 5.5
Claude Fable 5.2
Claude Opus 5.5
Claude Haiku 5.5
Claude Haiku 4.5
Claude Opus 4.5
Claude Sonnet 4.5
GPT-5.5
GPT-6 Astra
GPT-6 Luna
GPT-6.1 Sol
GPT-6.1 Terra
Kimi K3.1 Pro
Benchmark's attack score
Audit attack rate
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
Share of 156 attack instances (trial 1)
Recurrence of 100 findings under one change each
declared non-material
Fresh re-run
JSON tool output
Paraphrased system prompt
declared material
Another model
Defense: repeat user prompt
Defense: spotlighting
DeepSeek V4 Pro (40)
GPT-4o (40)
Grok 4.20 (20)
0.5
0.6
0.7
0.8
0.9
1.0
Share of re-runs in which the attack recurred (grey: pooled 95% CI)