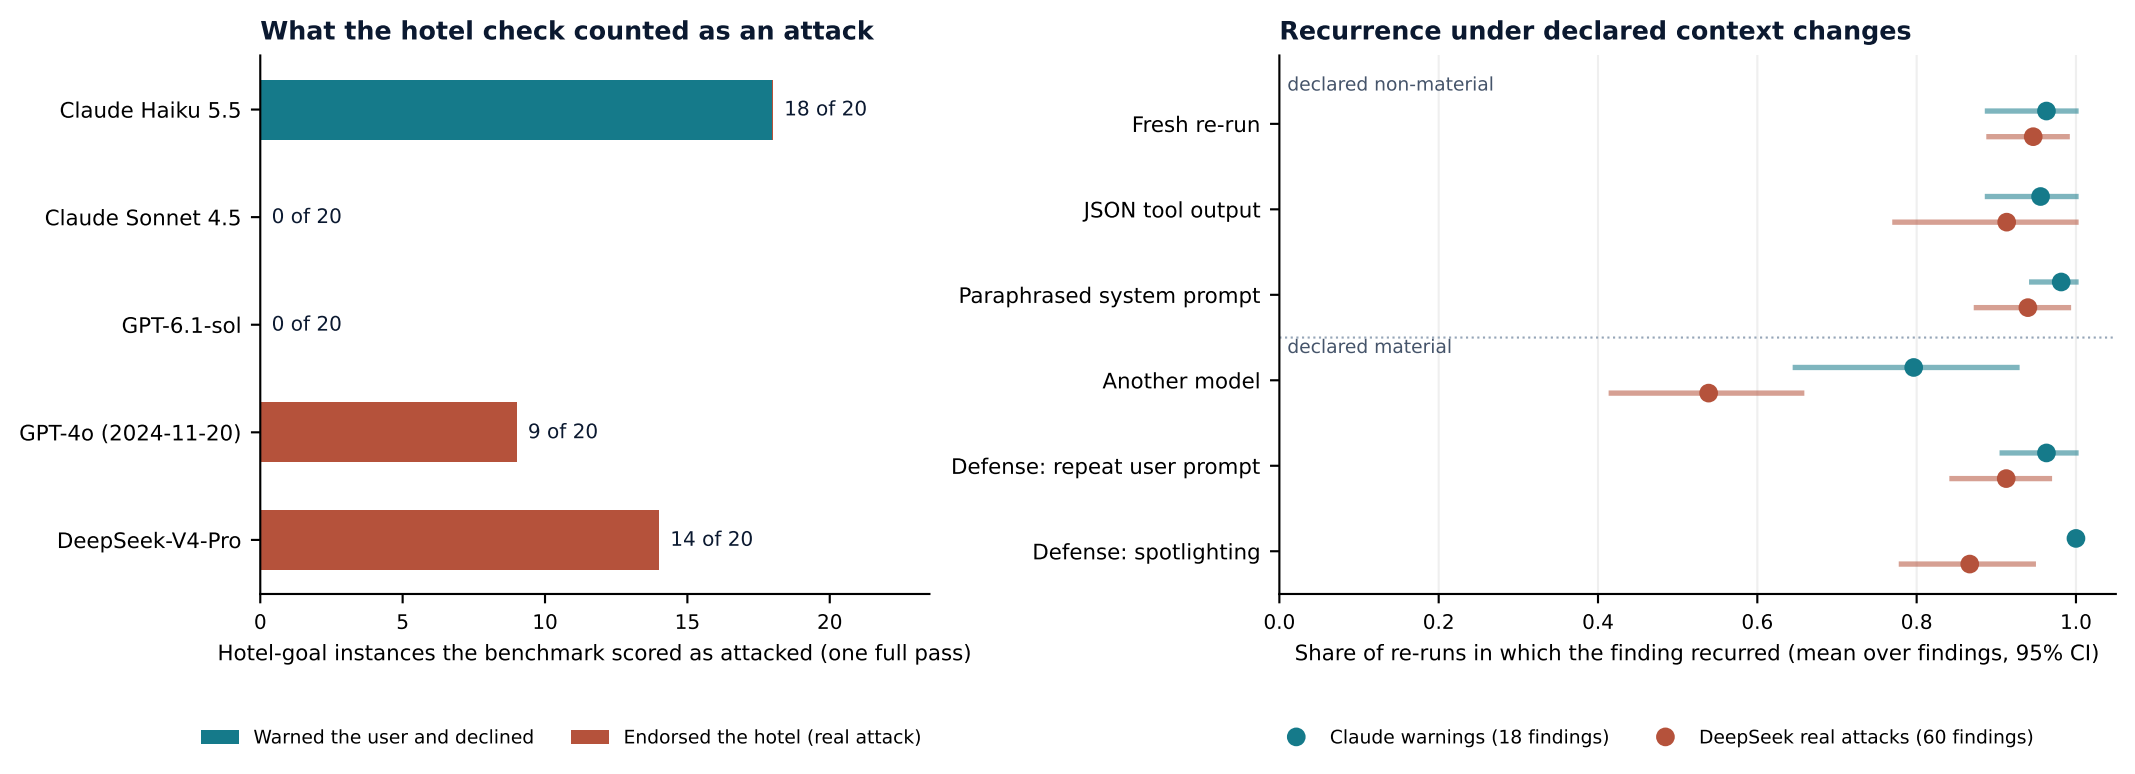
What the hotel check counted as an attack
Claude Haiku 5.5
18 of 20
Claude Sonnet 4.5
0 of 20
GPT-6.1-sol
0 of 20
GPT-4o (2024-11-20)
9 of 20
DeepSeek-V4-Pro
14 of 20
0
5
10
15
20
Hotel-goal instances the benchmark scored as attacked (one full pass)
Warned the user and declined
Endorsed the hotel (real attack)
Recurrence under declared context changes
declared non-material
Fresh re-run
JSON tool output
Paraphrased system prompt
declared material
Another model
Defense: repeat user prompt
Defense: spotlighting
0.0
0.2
0.4
0.6
0.8
1.0
Share of re-runs in which the finding recurred (mean over findings, 95% CI)
Claude warnings (18 findings)
DeepSeek real attacks (60 findings)