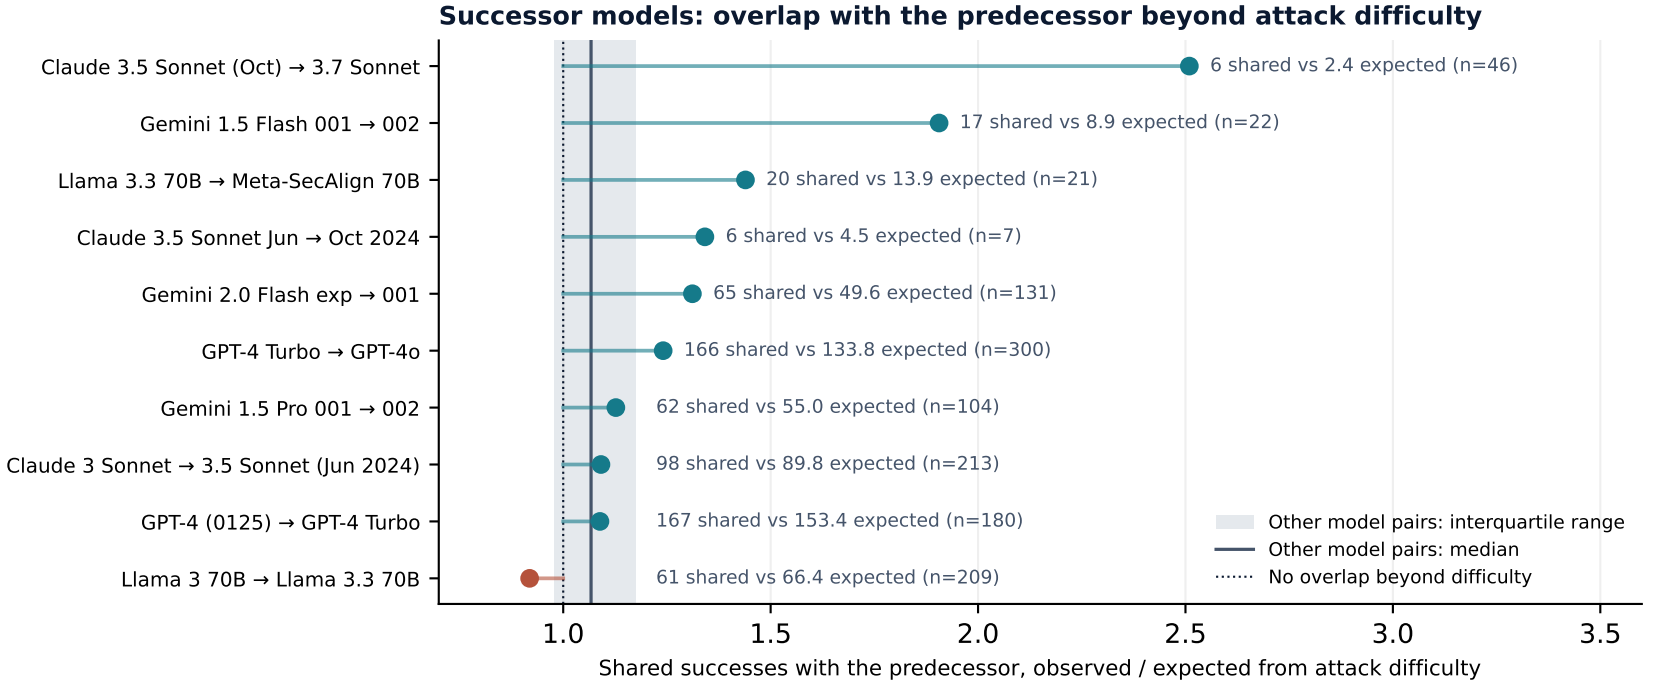
Successor models: overlap with the predecessor beyond attack difficulty
Claude 3.5 Sonnet (Oct) → 3.7 Sonnet
6 shared vs 2.4 expected (n=46)
Gemini 1.5 Flash 001 → 002
17 shared vs 8.9 expected (n=22)
Llama 3.3 70B → Meta-SecAlign 70B
20 shared vs 13.9 expected (n=21)
Claude 3.5 Sonnet Jun → Oct 2024
6 shared vs 4.5 expected (n=7)
Gemini 2.0 Flash exp → 001
65 shared vs 49.6 expected (n=131)
GPT-4 Turbo → GPT-4o
166 shared vs 133.8 expected (n=300)
Gemini 1.5 Pro 001 → 002
62 shared vs 55.0 expected (n=104)
Claude 3 Sonnet → 3.5 Sonnet (Jun 2024)
98 shared vs 89.8 expected (n=213)
GPT-4 (0125) → GPT-4 Turbo
167 shared vs 153.4 expected (n=180)
Llama 3 70B → Llama 3.3 70B
61 shared vs 66.4 expected (n=209)
Other model pairs: interquartile range
Other model pairs: median
No overlap beyond difficulty
1.0
1.5
2.0
2.5
3.0
3.5
Shared successes with the predecessor, observed / expected from attack difficulty